

Task 1

```
In [58]: #Import All the Necessary Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib as mpl
from mpl_toolkits.mplot3d import Axes3D
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler,QuantileTransformer
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

%matplotlib inline
```

Working on Train dataframe

```
In [59]: traindf = pd.read_csv('train.csv')
```

```
In [3]: traindf.columns
```

```
Out[3]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
              'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
              'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
              'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
              'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
              'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
              'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
              'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
              'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
              'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
              'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
              'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
              'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
              'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
              'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
              'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
              'SaleCondition', 'SalePrice'],
              dtype='object')
```

```
In [4]: numeric_df = traindf.select_dtypes(include='number')
correlation_matrix = numeric_df.corr()
correlation_matrix['SalePrice'].sort_values(ascending = False)
```

```
Out[4]: SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
TotalBsmstSF   0.613581
1stFlrSF       0.605852
FullBath       0.560664
TotRmsAbvGrd   0.533723
YearBuilt      0.522897
YearRemodAdd    0.507101
GarageYrBlt    0.486362
MasVnrArea     0.477493
Fireplaces     0.466929
BsmstFinSF1    0.386420
LotFrontage    0.351799
WoodDeckSF     0.324413
2ndFlrSF       0.319334
OpenPorchSF    0.315856
HalfBath       0.284108
LotArea        0.263843
BsmstFullBath  0.227122
BsmstUnfSF     0.214479
BedroomAbvGr   0.168213
ScreenPorch    0.111447
PoolArea       0.092404
MoSold         0.046432
3SsnPorch      0.044584
BsmstFinSF2    -0.011378
BsmstHalfBath  -0.016844
MiscVal        -0.021190
Id             -0.021917
LowQualFinSF   -0.025606
YrSold         -0.028923
OverallCond    -0.077856
MSSubClass     -0.084284
EnclosedPorch  -0.128578
KitchenAbvGr   -0.135907
Name: SalePrice, dtype: float64
```

```
In [5]: req_tr = ["GarageArea", "OverallQual", "TotalBsmstSF", "1stFlrSF", "2ndFlrSF", "LowQualFinSF", "GrLivArea", "BsmstFullBath", "BsmstHalfBath"]
```

```
In [6]: selected_tr = traindf[req_tr]
```

```
In [12]: selected_tr.loc[:, 'TotalBath'] = (selected_tr['BsmstFullBath'].fillna(0) +
                                             selected_tr['BsmstHalfBath'].fillna(0) +
                                             selected_tr['FullBath'].fillna(0) +
                                             selected_tr['HalfBath'].fillna(0))

selected_tr.loc[:, 'TotalSF'] = (selected_tr['TotalBsmstSF'].fillna(0) +
                                  selected_tr['1stFlrSF'].fillna(0) +
                                  selected_tr['2ndFlrSF'].fillna(0) +
                                  selected_tr['LowQualFinSF'].fillna(0) +
                                  selected_tr['GrLivArea'].fillna(0))
```

In [8]: selected_tr

Out[8]:

	GarageArea	OverallQual	TotalBsmtSF	1stFlrSF	2ndFlrSF	LowQualFinSF	GrLivArea	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	TotRmsAbvGrd	S
0	548	7	856	856	854	0	1710	1	0	2	1	8	
1	460	6	1262	1262	0	0	1262	0	1	2	0	6	
2	608	7	920	920	866	0	1786	1	0	2	1	6	
3	642	7	756	961	756	0	1717	1	0	1	0	7	
4	836	8	1145	1145	1053	0	2198	1	0	2	1	9	
...	
1455	460	6	953	953	694	0	1647	0	0	2	1	7	
1456	500	6	1542	2073	0	0	2073	1	0	2	0	7	
1457	252	7	1152	1188	1152	0	2340	0	0	2	0	9	
1458	240	5	1078	1078	0	0	1078	1	0	1	0	5	
1459	276	5	1256	1256	0	0	1256	1	0	1	1	6	

1460 rows × 15 columns

Keeping only the necessary columns

In [9]: train_df = selected_tr[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual', 'SalePrice']]

In [10]: train_df

Out[10]:

	TotRmsAbvGrd	TotalBath	GarageArea	TotalSF	OverallQual	SalePrice
0	8	4	548	4276	7	208500
1	6	3	460	3786	6	181500
2	6	4	608	4492	7	223500
3	7	2	642	4190	7	140000
4	9	4	836	5541	8	250000
...
1455	7	3	460	4247	6	175000
1456	7	3	500	5688	6	210000
1457	9	2	252	5832	7	266500
1458	5	2	240	3234	5	142125
1459	6	3	276	3768	5	147500

1460 rows × 6 columns

Splitting the dataset and Creating Pipeline

```
In [14]: from sklearn.model_selection import train_test_split
train_set, test_set = train_test_split(train_df, test_size = 0.2, random_state = 42)
print(f"Rows in train set: {len(train_set)}\nRows in test set: {len(test_set)}\n")
```

Rows in train set: 1168
Rows in test set: 292

```
In [15]: housing = train_set.drop("SalePrice", axis=1)
housing_labels = train_set["SalePrice"].copy()
```

```
In [16]: from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
my_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler())
])
```

```
In [17]: X_train = my_pipeline.fit_transform(housing)
```

In [18]: X_train

Out[18]: array([[-0.96456591, -0.48377079, -0.86383727, -0.13352109, -0.82044456],
[0.27075534, 0.61127627, -0.45626397, -0.13428593, -0.08893368],
[-1.58222654, -1.57881784, -2.25716927, -1.32207838, -0.82044456],
...,
[-0.96456591, -0.48377079, 0.45366713, -1.16605156, -0.82044456],
[0.27075534, -0.48377079, -1.23349678, -0.26966215, 0.64257719],
[0.27075534, -0.48377079, 0.87071888, 0.28025593, 0.64257719]])

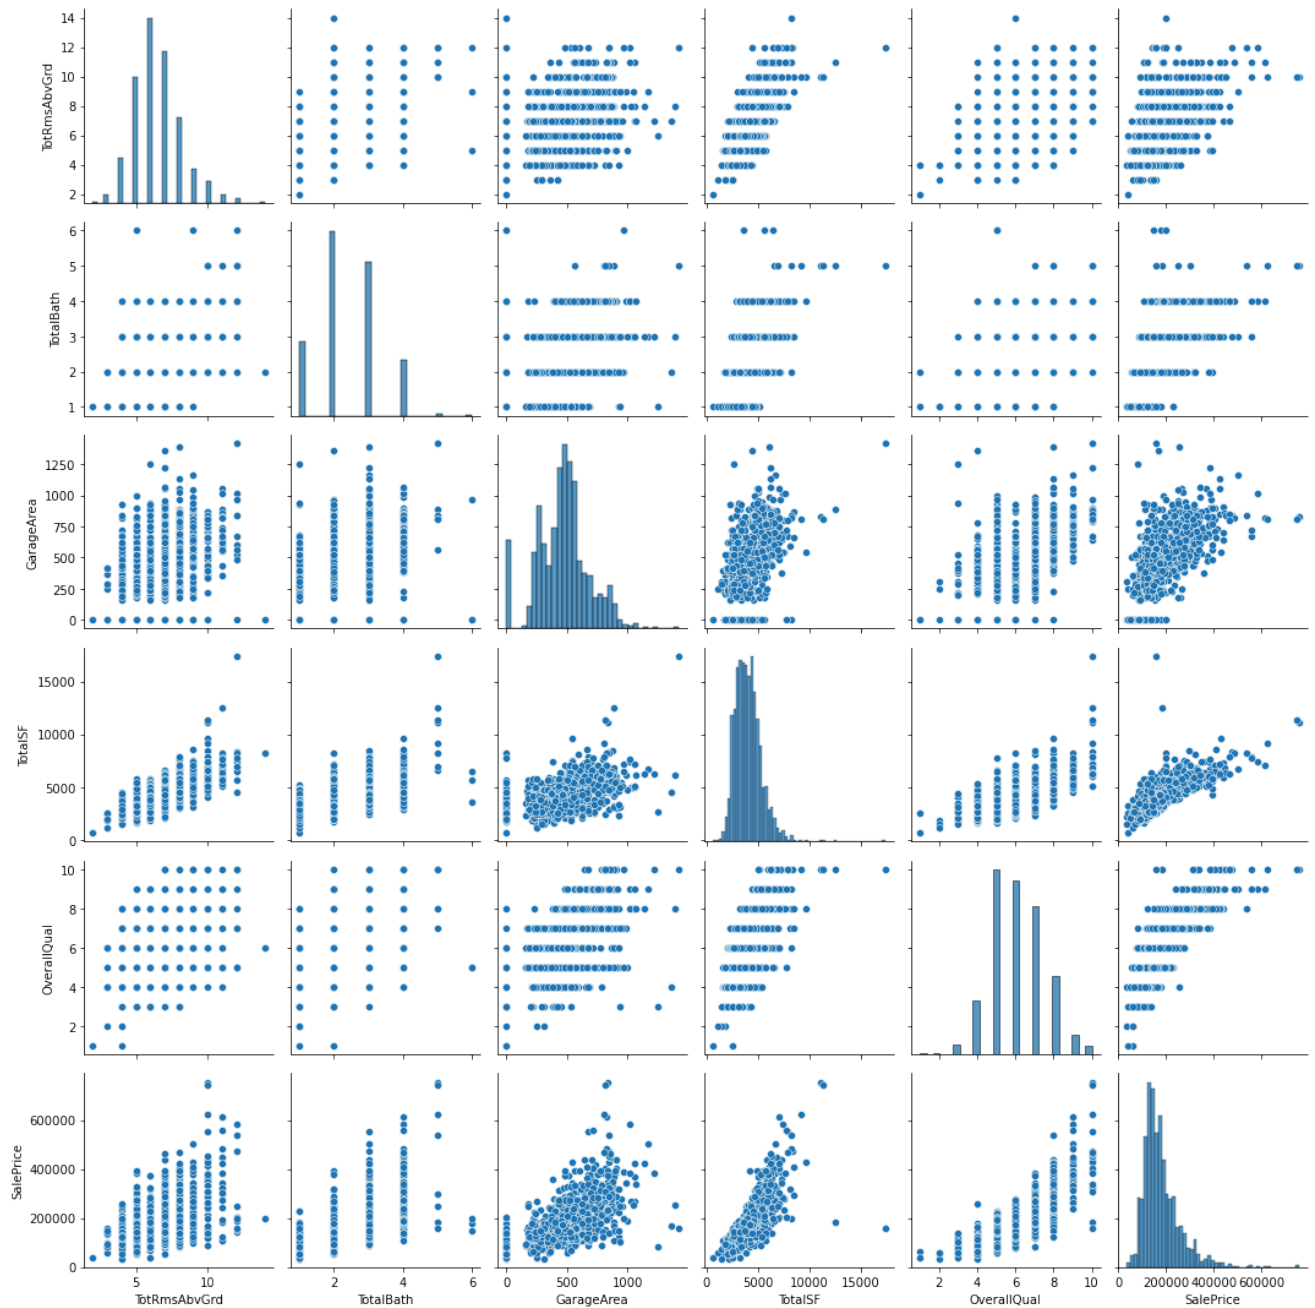
In [19]: Y_train = housing_labels

In [20]: Y_train.shape

Out[20]: (1168,)


Correlations

In [22]: `import warnings`
`warnings.filterwarnings("ignore", category=UserWarning)`
`%matplotlib inline`
`sns.pairplot(train_df)`
`plt.tight_layout()`
`plt.show()`



```
In [23]: sns.heatmap(train_df.corr(),annot = True)
```

```
Out[23]: <AxesSubplot:>
```



Working with Test Dataframe

```
In [25]: testdf = pd.read_csv("test.csv")
testdf.head()
```

```
Out[25]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence	MiscFeature
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllPub	...	120	0	NaN	MnPrv	Na
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllPub	...	0	0	NaN	NaN	Ga
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllPub	...	0	0	NaN	MnPrv	Na
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	Lvl	AllPub	...	0	0	NaN	NaN	Na
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	HLS	AllPub	...	144	0	NaN	NaN	Na

5 rows × 80 columns

```
In [26]: req_tst = ["GarageArea", "OverallQual", "TotalBsmtSF", "1stFlrSF", "2ndFlrSF", "LowQualFinSF", "GrLivArea", "BsmtFullBath", "BsmtHalfBath"]
selected_tst = testdf[req_tst]
```

```
In [27]: selected_tst = testdf[req_tst]
```

```
In [32]: selected_tst.loc[:, 'TotalBath'] = (selected_tst['BsmtFullBath'].fillna(0) +
                                             selected_tst['BsmtHalfBath'].fillna(0) +
                                             selected_tst['FullBath'].fillna(0) +
                                             selected_tst['HalfBath'].fillna(0))

selected_tst.loc[:, 'TotalSF'] = (selected_tst['TotalBsmtSF'].fillna(0) +
                                   selected_tst['1stFlrSF'].fillna(0) +
                                   selected_tst['2ndFlrSF'].fillna(0) +
                                   selected_tst['LowQualFinSF'].fillna(0) +
                                   selected_tst['GrLivArea'].fillna(0))
```

In [30]: selected_tst

Out[30]:

	GarageArea	OverallQual	TotalBsmtSF	1stFlrSF	2ndFlrSF	LowQualFinSF	GrLivArea	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	TotRmsAbvGrd	T
0	730.0	5	882.0	896	0	0	896	0.0	0.0	1	0	5	
1	312.0	6	1329.0	1329	0	0	1329	0.0	0.0	1	1	6	
2	482.0	5	928.0	928	701	0	1629	0.0	0.0	2	1	6	
3	470.0	6	926.0	926	678	0	1604	0.0	0.0	2	1	7	
4	506.0	8	1280.0	1280	0	0	1280	0.0	0.0	2	0	5	
...	
1454	0.0	4	546.0	546	546	0	1092	0.0	0.0	1	1	5	
1455	286.0	4	546.0	546	546	0	1092	0.0	0.0	1	1	6	
1456	576.0	5	1224.0	1224	0	0	1224	1.0	0.0	1	0	7	
1457	0.0	5	912.0	970	0	0	970	0.0	1.0	1	0	6	
1458	650.0	7	996.0	996	1004	0	2000	0.0	0.0	2	1	9	

1459 rows × 14 columns

In [33]: test_df_unproc = selected_tst[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual']]

In [34]: test_df_unproc

Out[34]:

	TotRmsAbvGrd	TotalBath	GarageArea	TotalSF	OverallQual
0	5	1.0	730.0	2674.0	5
1	6	2.0	312.0	3987.0	6
2	6	3.0	482.0	4186.0	5
3	7	3.0	470.0	4134.0	6
4	5	2.0	506.0	3840.0	8
...
1454	5	2.0	0.0	2730.0	4
1455	6	2.0	286.0	2730.0	4
1456	7	2.0	576.0	3672.0	5
1457	6	2.0	0.0	2852.0	5
1458	9	3.0	650.0	4996.0	7

1459 rows × 5 columns

In [35]: test_df = test_df_unproc.fillna(test_df_unproc.mean())

In [36]: x_test = my_pipeline.transform(test_df[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual']].values)

In [37]: x_test

Out[37]: array([[-0.96456591, -1.57881784, 1.2024646 , -1.10333489, -0.82044456],
 [-0.34690528, -0.48377079, -0.77853123, -0.09910341, -0.08893368],
 [-0.34690528, 0.61127627, 0.02713693, 0.05309923, -0.82044456],
 ...,
 [0.27075534, -0.48377079, 0.47262403, -0.34002719, -0.82044456],
 [-0.34690528, -0.48377079, -2.25716927, -0.96719384, -0.82044456],
 [1.50607659, 0.61127627, 0.82332664, 0.67261751, 0.64257719]])

Model Selection

In [39]: #model = LinearRegression()
 #model = DecisionTreeRegressor()
 model = RandomForestRegressor()
 model.fit(X_train, Y_train)

Out[39]: RandomForestRegressor()

In [40]: y_train_pred = model.predict(X_train)

```
In [41]: y_train_pred[:5]
```

```
Out[41]: array([146834.84, 172931.32,  91300.  , 166744.24, 140191.  ])
```

```
In [42]: some_data = housing.iloc[:5]  
some_labels = housing_labels.iloc[:5]
```

```
In [43]: proc_data = my_pipeline.transform(some_data)
```

```
In [44]: model.predict(proc_data)
```

```
Out[44]: array([146834.84, 172931.32,  91300.  , 166744.24, 140191.  ])
```

```
In [45]: list(some_labels)
```

```
Out[45]: [145000, 178000, 85000, 175000, 127000]
```

```
In [46]: train_mse = mean_squared_error(Y_train,y_train_pred)
```

```
In [47]: train_rmse = np.sqrt(train_mse)
```

```
In [48]: print(f"Training MSE: {train_mse:.2f}, Training RMSE: {train_rmse:.2f}")
```

```
Training MSE: 172026550.98, Training RMSE: 13115.89
```

```
Cross - Validation
```

```
In [50]: from sklearn.model_selection import cross_val_score  
scores = cross_val_score(model,X_train,Y_train,scoring="neg_mean_squared_error",cv = 200)  
rmse_scores = np.sqrt(-scores)
```

In [52]: rmse_scores

```
Out[52]: array([[ 21797.00669202, 14957.23898344, 23275.47683597, 11572.31366899,
 48469.42139054, 7978.34457094, 20108.00276321, 12182.28705366,
 9675.6005755 , 50051.95097995, 35336.34359263, 28413.51341703,
 13376.37160055, 9231.85991138, 19188.14282691, 24133.7195857 ,
 22196.33020079, 34130.85574345, 37975.93581947, 22577.63227323,
 29179.25399623, 18375.66206345, 17160.70071119, 27878.48541779,
 19603.82459691, 16825.41801781, 43740.14145525, 39724.65220952,
 176909.66719632, 51788.88255228, 23334.31657806, 33784.70364157,
 22282.38778279, 32334.5193015 , 50705.24900549, 13047.87527871,
 22625.86361465, 29734.56486361, 19631.81205684, 30324.90502457,
 21632.85519688, 31509.02307562, 28060.15770012, 32925.9132895 ,
 33147.81554893, 29222.27202686, 27771.70311292, 42393.55519823,
 21369.22916065, 22467.10057453, 20924.18662807, 48161.4827102 ,
 41136.68890422, 36070.41333717, 23224.24710799, 26647.34160335,
 10923.56368389, 27395.41929729, 25133.96265023, 30794.43284436,
 200176.50686324, 12964.3965824 , 23573.53115501, 35304.32045918,
 22433.83451002, 25678.94236018, 37458.0628356 , 29555.53944044,
 7661.74250814, 14727.787973 , 57329.12632426, 27221.77836773,
 57520.85344648, 18551.85315539, 46763.06902435, 43296.50993689,
 40860.09435628, 29802.20027902, 41610.13598278, 32778.20515812,
 27980.9459166 , 32012.30602063, 102727.14956704, 8512.05566991,
 48702.8238754 , 33746.41055608, 18892.61485873, 23173.03608053,
 34232.51683334, 76160.87323065, 19534.11024824, 22763.47101524,
 26778.91449607, 24839.86969716, 22023.51121385, 22899.09792426,
 25177.27969627, 27054.33547047, 70353.74989691, 13075.34584013,
 30295.02368228, 27852.836771 , 47272.47658452, 52837.22871236,
 18265.70827883, 18609.07531439, 37755.77969763, 18381.18261301,
 19391.26991251, 15537.5580907 , 19831.66264374, 12679.66154048,
 12374.18002601, 19689.58169192, 29660.72321625, 34117.9541451 ,
 76540.12052656, 29696.42354019, 18924.90446707, 25584.9575775 ,
 30115.01267773, 28243.26709808, 16246.69567596, 20069.33052675,
 33494.84124277, 20665.78796561, 30613.53973346, 46517.14561505,
 39252.64005476, 16479.40451005, 33323.31507917, 7912.69298667,
 24606.93935367, 27606.29374646, 26867.35518677, 12344.08722316,
 47467.0233914 , 34801.16337931, 28816.48668026, 12774.76565494,
 25667.81631088, 25781.74953679, 26568.81983178, 11793.86336533,
 27023.99016464, 28976.83715756, 32867.98942372, 13969.1685425 ,
 12431.51201239, 15682.35086843, 25441.42785125, 26944.52144272,
 14107.63519776, 41800.21756193, 30936.17488118, 11967.64690939,
 23653.52173679, 25313.56984546, 21167.41561914, 23787.55509279,
 52678.95257977, 27239.65401155, 26241.70676363, 19016.82308447,
 49723.35586252, 21594.08841625, 12986.52271444, 61732.56257263,
 20709.83670122, 29636.14876585, 29617.10636478, 23944.94913123,
 14187.16088132, 13210.82900748, 14164.14058122, 18636.91790398,
 22962.13911521, 16815.15778132, 39186.89137075, 21304.89640533,
 33535.667977 , 11000.74891384, 15984.82892073, 6314.40760104,
 17886.30632806, 31777.16028452, 26004.08548421, 13937.61202777,
 26535.04161049, 38994.06371719, 29611.16999628, 19890.36672566,
 17924.8571899 , 13186.03558123, 15470.62428205, 17103.37379449,
 58208.00583386, 20068.74696506, 17829.19613253, 29203.44959104])
```

```
In [53]: def print_scores(scores):
          print("Scores:", scores)
          print("Mean:", scores.mean())
          print("Standard Deviation", scores.std())
```


In [54]: `print_scores(rmse_scores)`

```
Scores: [ 21797.00669202 14957.23898344 23275.47683597 11572.31366899
 48469.42139054  7978.34457094 20108.00276321 12182.28705366
 9675.6005755  50051.95097995 35336.34359263 28413.51341703
13376.37160055  9231.85991138 19188.14282691 24133.7195857
22196.33020079 34130.85574345 37975.93581947 22577.63227323
29179.25399623 18375.66206345 17160.70071119 27878.48541779
19603.82459691 16825.41801781 43740.14145525 39724.65220952
176909.66719632 51788.88255228 23334.31657806 33784.70364157
22282.38778279 32334.5193015 50705.24900549 13047.87527871
22625.86361465 29734.56486361 19631.81205684 30324.90502457
21632.85519688 31509.02307562 28060.15770012 32925.9132895
33147.81554893 29222.27202686 27771.70311292 42393.55519823
21369.22916065 22467.10057453 20924.18662807 48161.4827102
41136.68890422 36070.41333717 23224.24710799 26647.34160335
10923.56368389 27395.41929729 25133.96265023 30794.43284436
200176.50686324 12964.3965824 23573.53115501 35304.32045918
22433.83451002 25678.94236018 37458.0628356 29555.53944044
7661.74250814 14727.787973 57329.12632426 27221.77836773
57520.85344648 18551.85315539 46763.06902435 43296.50993689
40860.09435628 29802.20027902 41610.13598278 32778.20515812
27980.9459166 32012.30602063 102727.14956704 8512.05566991
48702.8238754 33746.41055608 18892.61485873 23173.03608053
34232.51683334 76160.87323065 19534.11024824 22763.47101524
26778.91449607 24839.86969716 22023.51121385 22899.09792426
25177.27969627 27054.33547047 70353.74989691 13075.34584013
30295.02368228 27852.836771 47272.47658452 52837.22871236
18265.70827883 18609.07531439 37755.77969763 18381.18261301
19391.26991251 15537.5580907 19831.66264374 12679.66154048
12374.18002601 19689.58169192 29660.72321625 34117.9541451
76540.12052656 29696.42354019 18924.90446707 25584.9575775
30115.01267773 28243.26709808 16246.69567596 20069.33052675
33494.84124277 20665.78796561 30613.53973346 46517.14561505
39252.64005476 16479.40451005 33323.31507917 7912.69298667
24606.93935367 27606.29374646 26867.35518677 12344.08722316
47467.0233914 34801.16337931 28816.48668026 12774.76565494
25667.81631088 25781.74953679 26568.81983178 11793.86336533
27023.99016464 28976.83715756 32867.98942372 13969.1685425
12431.51201239 15682.35086843 25441.42785125 26944.52144272
14107.63519776 41800.21756193 30936.17488118 11967.64690939
23653.52173679 25313.56984546 21167.41561914 23787.55509279
52678.95257977 27239.65401155 26241.70676363 19016.82308447
49723.35586252 21594.08841625 12986.52271444 61732.56257263
20709.83670122 29636.14876585 29617.10636478 23944.94913123
14187.16088132 13210.82900748 14164.14058122 18636.91790398
22962.13911521 16815.15778132 39186.89137075 21304.89640533
33535.667977 11000.74891384 15984.82892073 6314.40760104
17886.30632806 31777.16028452 26004.08548421 13937.61202777
26535.04161049 38994.06371719 29611.16999628 19890.36672566
17924.8571899 13186.03558123 15470.62428205 17103.37379449
58208.00583386 20068.74696506 17829.19613253 29203.44959104]
Mean: 29089.48487732674
Standard Deviation 21087.336412413624
```

In [55]: `y_pred=model.predict(x_test)`

In [56]: `y_pred`

Out[56]: `array([133177.5 , 154364.32, 144508.37, ..., 140685. , 108002. ,
 235164.1])`

In [57]: `pred=pd.DataFrame(y_pred)
sub_df=pd.read_csv('sample_submission.csv')
datasets=pd.concat([sub_df['Id'],pred],axis=1)
datasets.columns=['Id','SalePrice']
datasets.to_csv('sample_submission.csv',index=False)`

In []: